


AWS Builders Online Series

Bring AWS to you with hybrid cloud storage

Naim Mucaj

Senior Solutions Architect
Amazon Web Services



Agenda

- Connectivity to AWS
- AWS Hybrid cloud storage services
- Common customer use cases and patterns
- Demo
- Resources



How much data is out there?

175 ZB

Zettabytes of data created
by 2025, globally

10^{21}

1 zettabyte is
1,000,000,000,000,000,000,000

Sources:

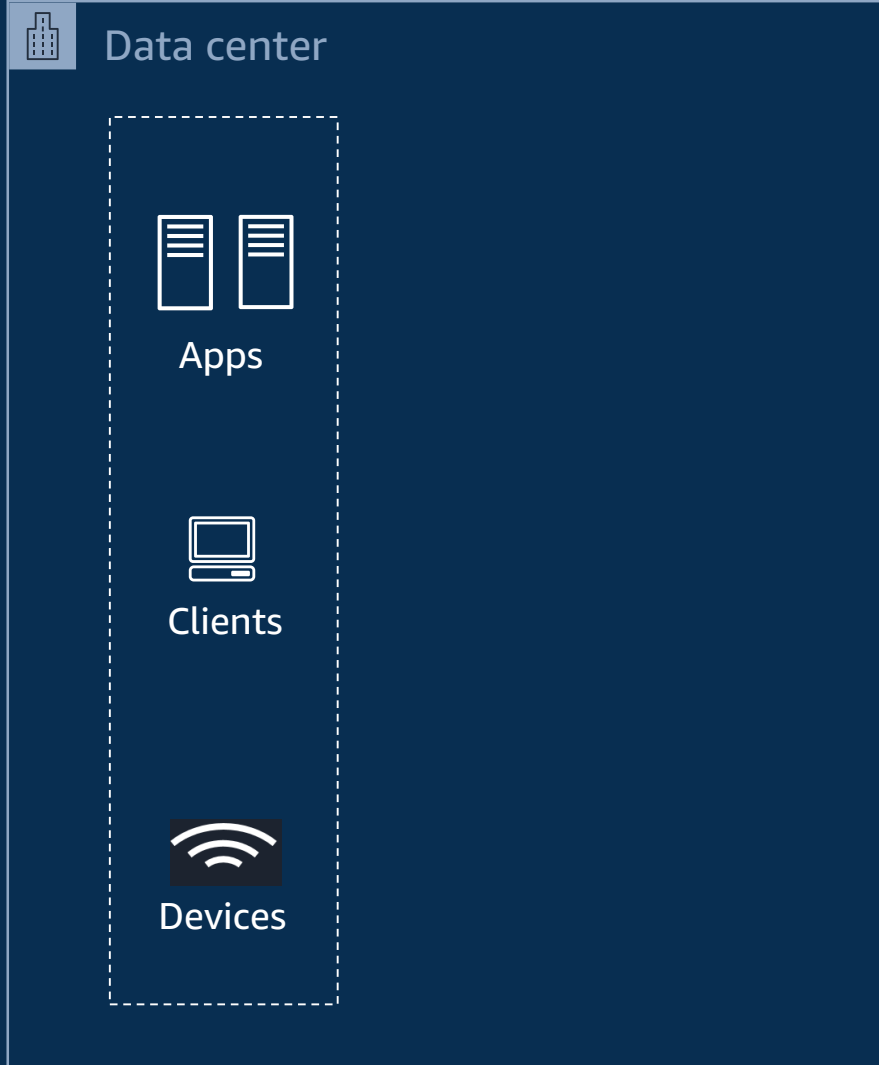
<https://www.idc.com/getdoc.jsp?containerId=prUS48247121>

<https://www.idc.com/getdoc.jsp?containerId=US47397521>


Network connectivity to AWS



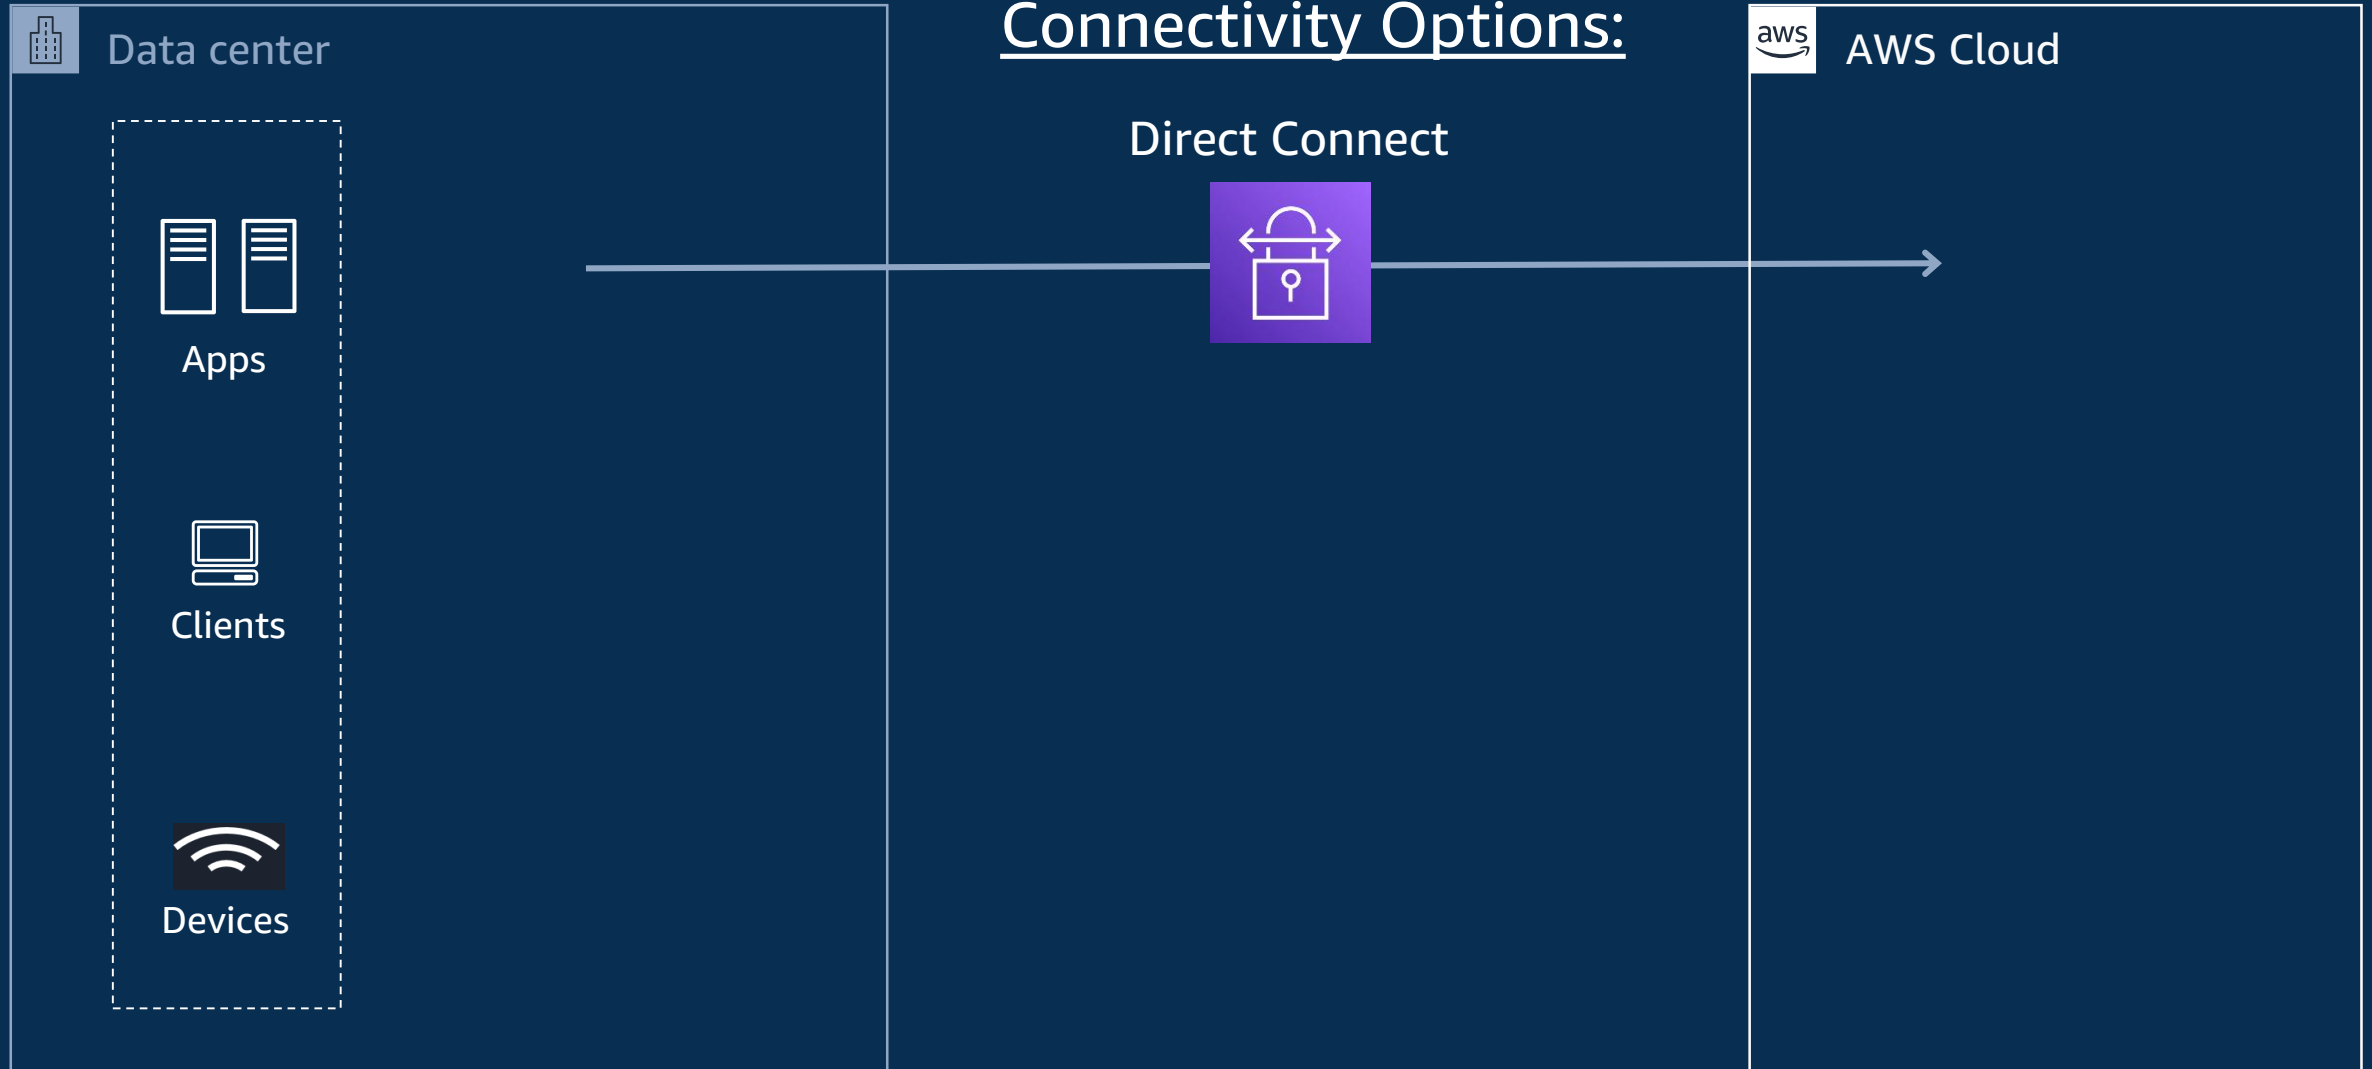
Connecting to AWS



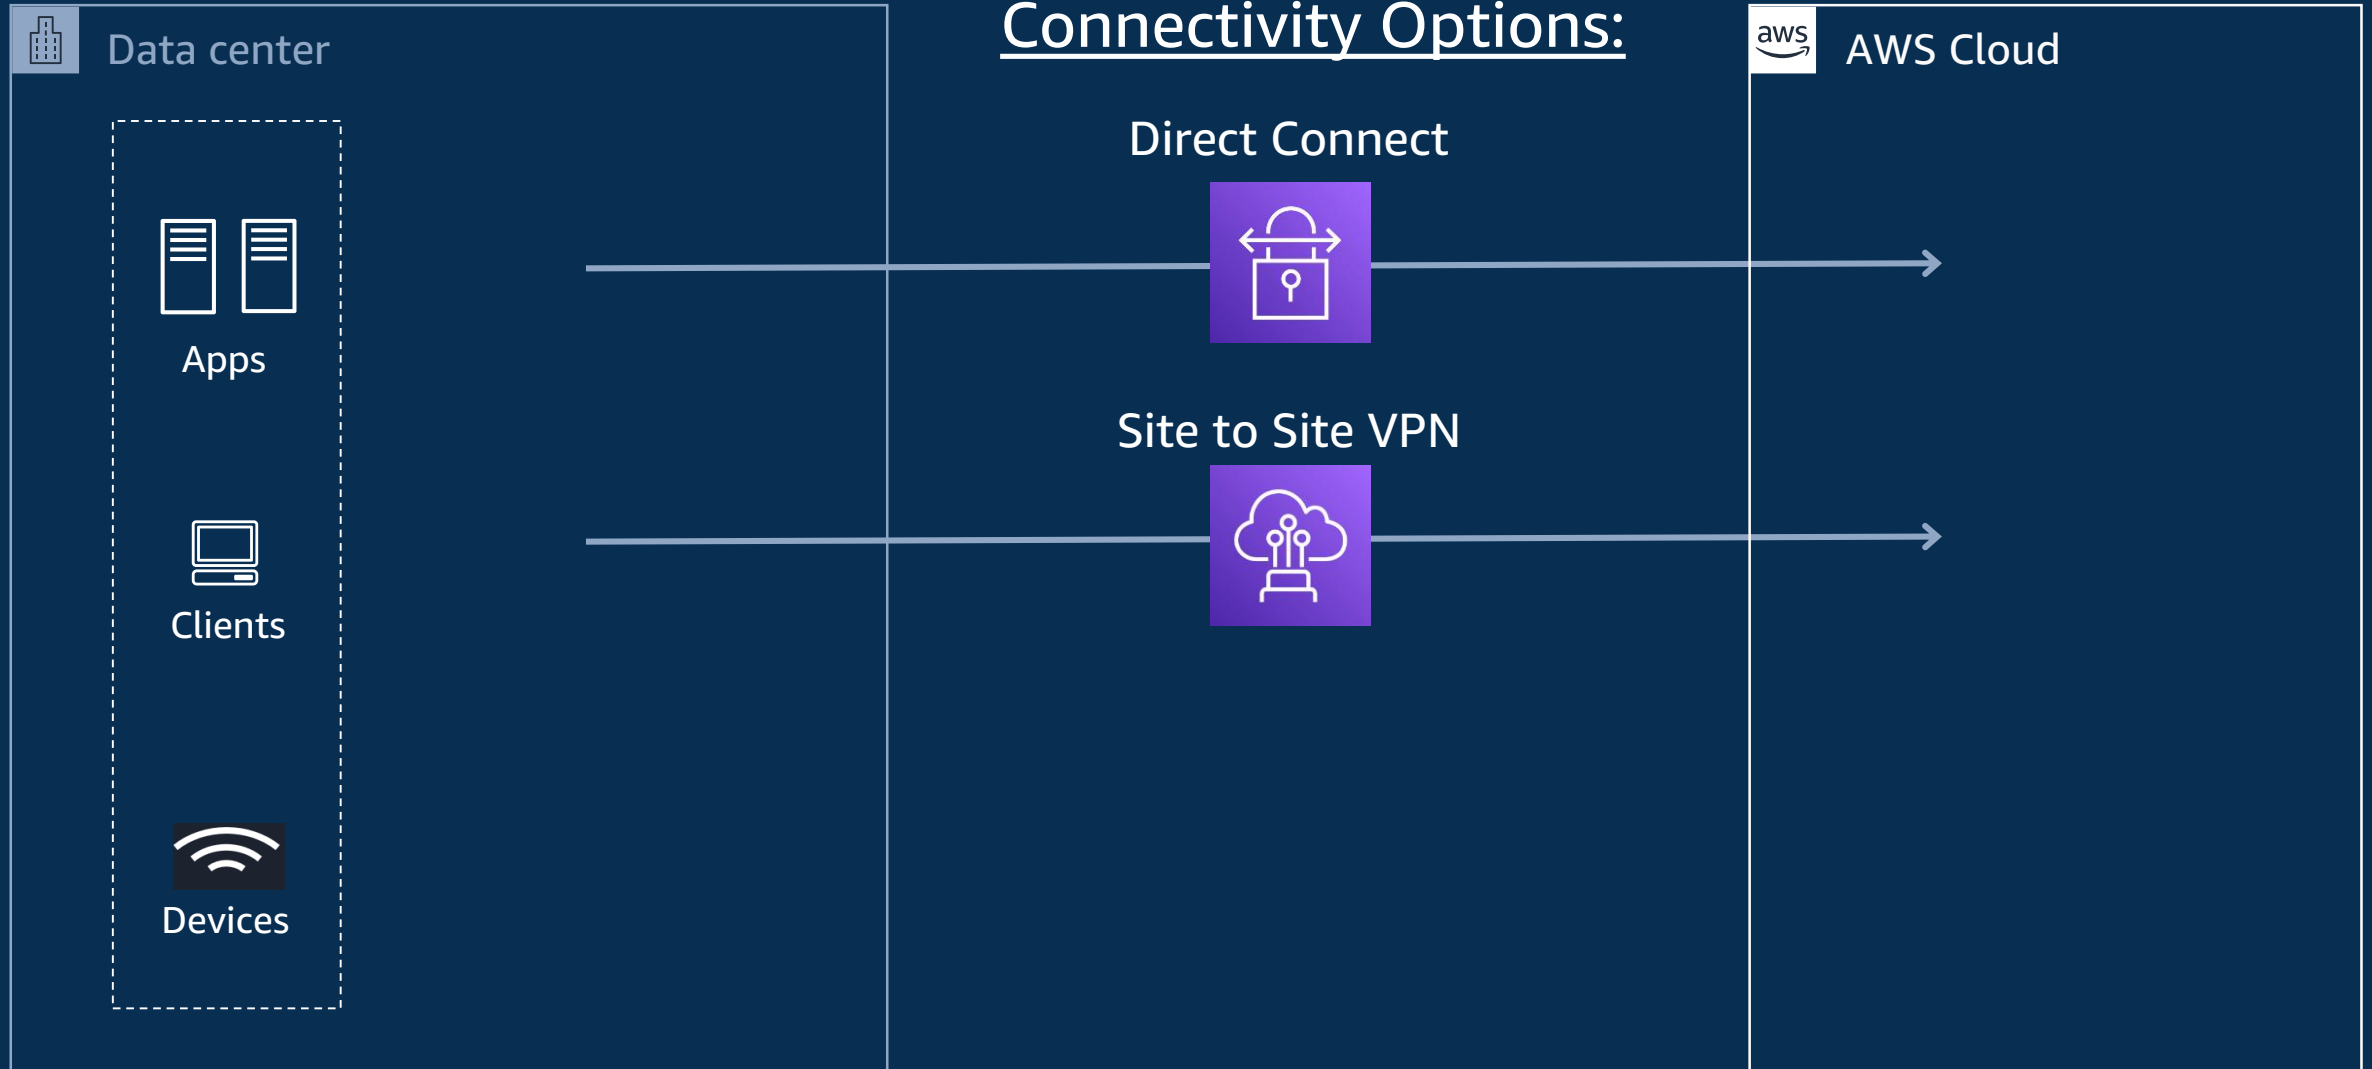
Connectivity Options:



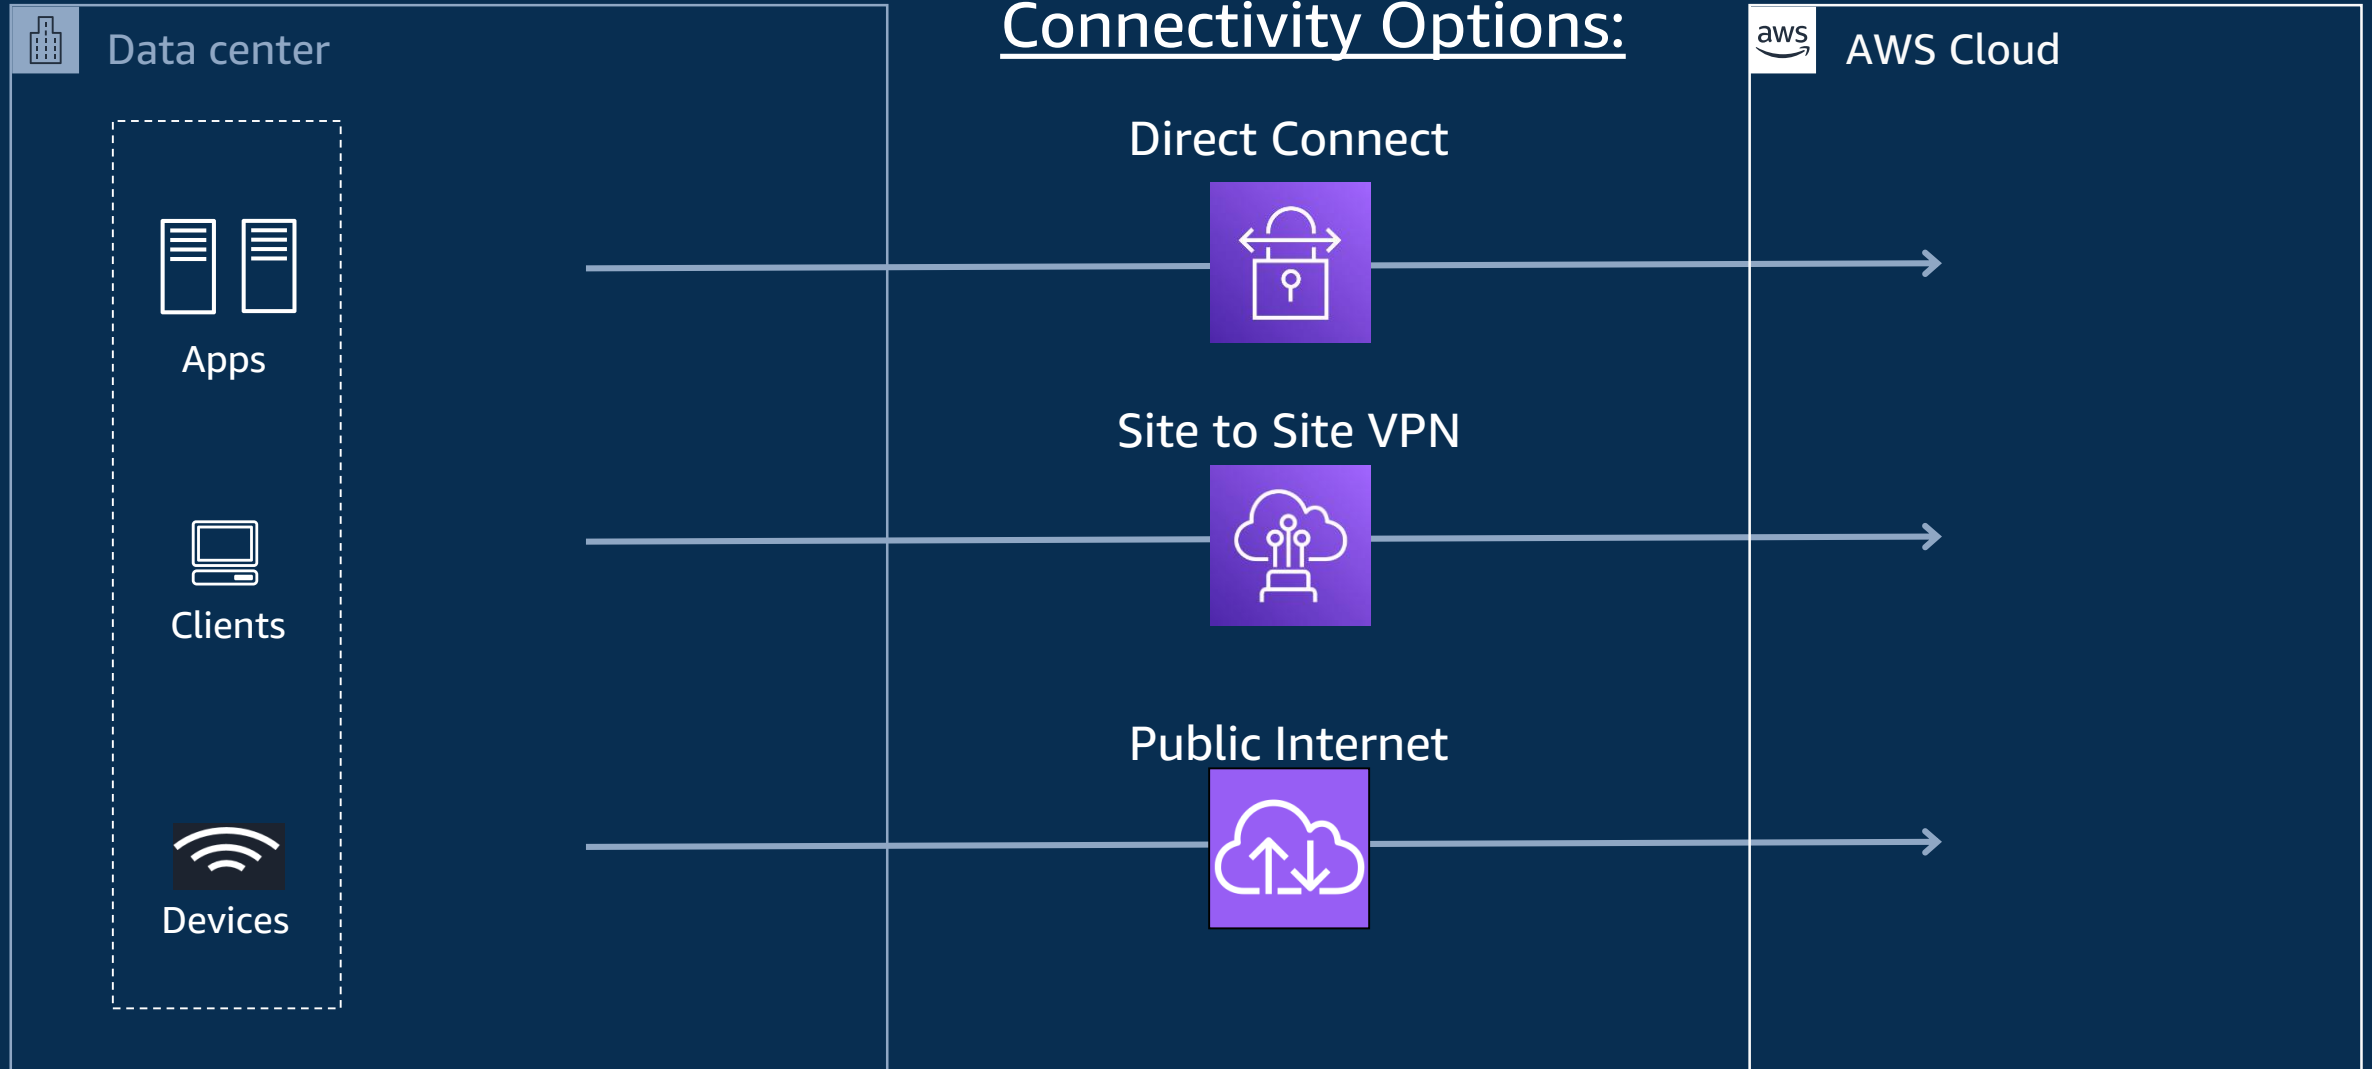
Connecting to AWS



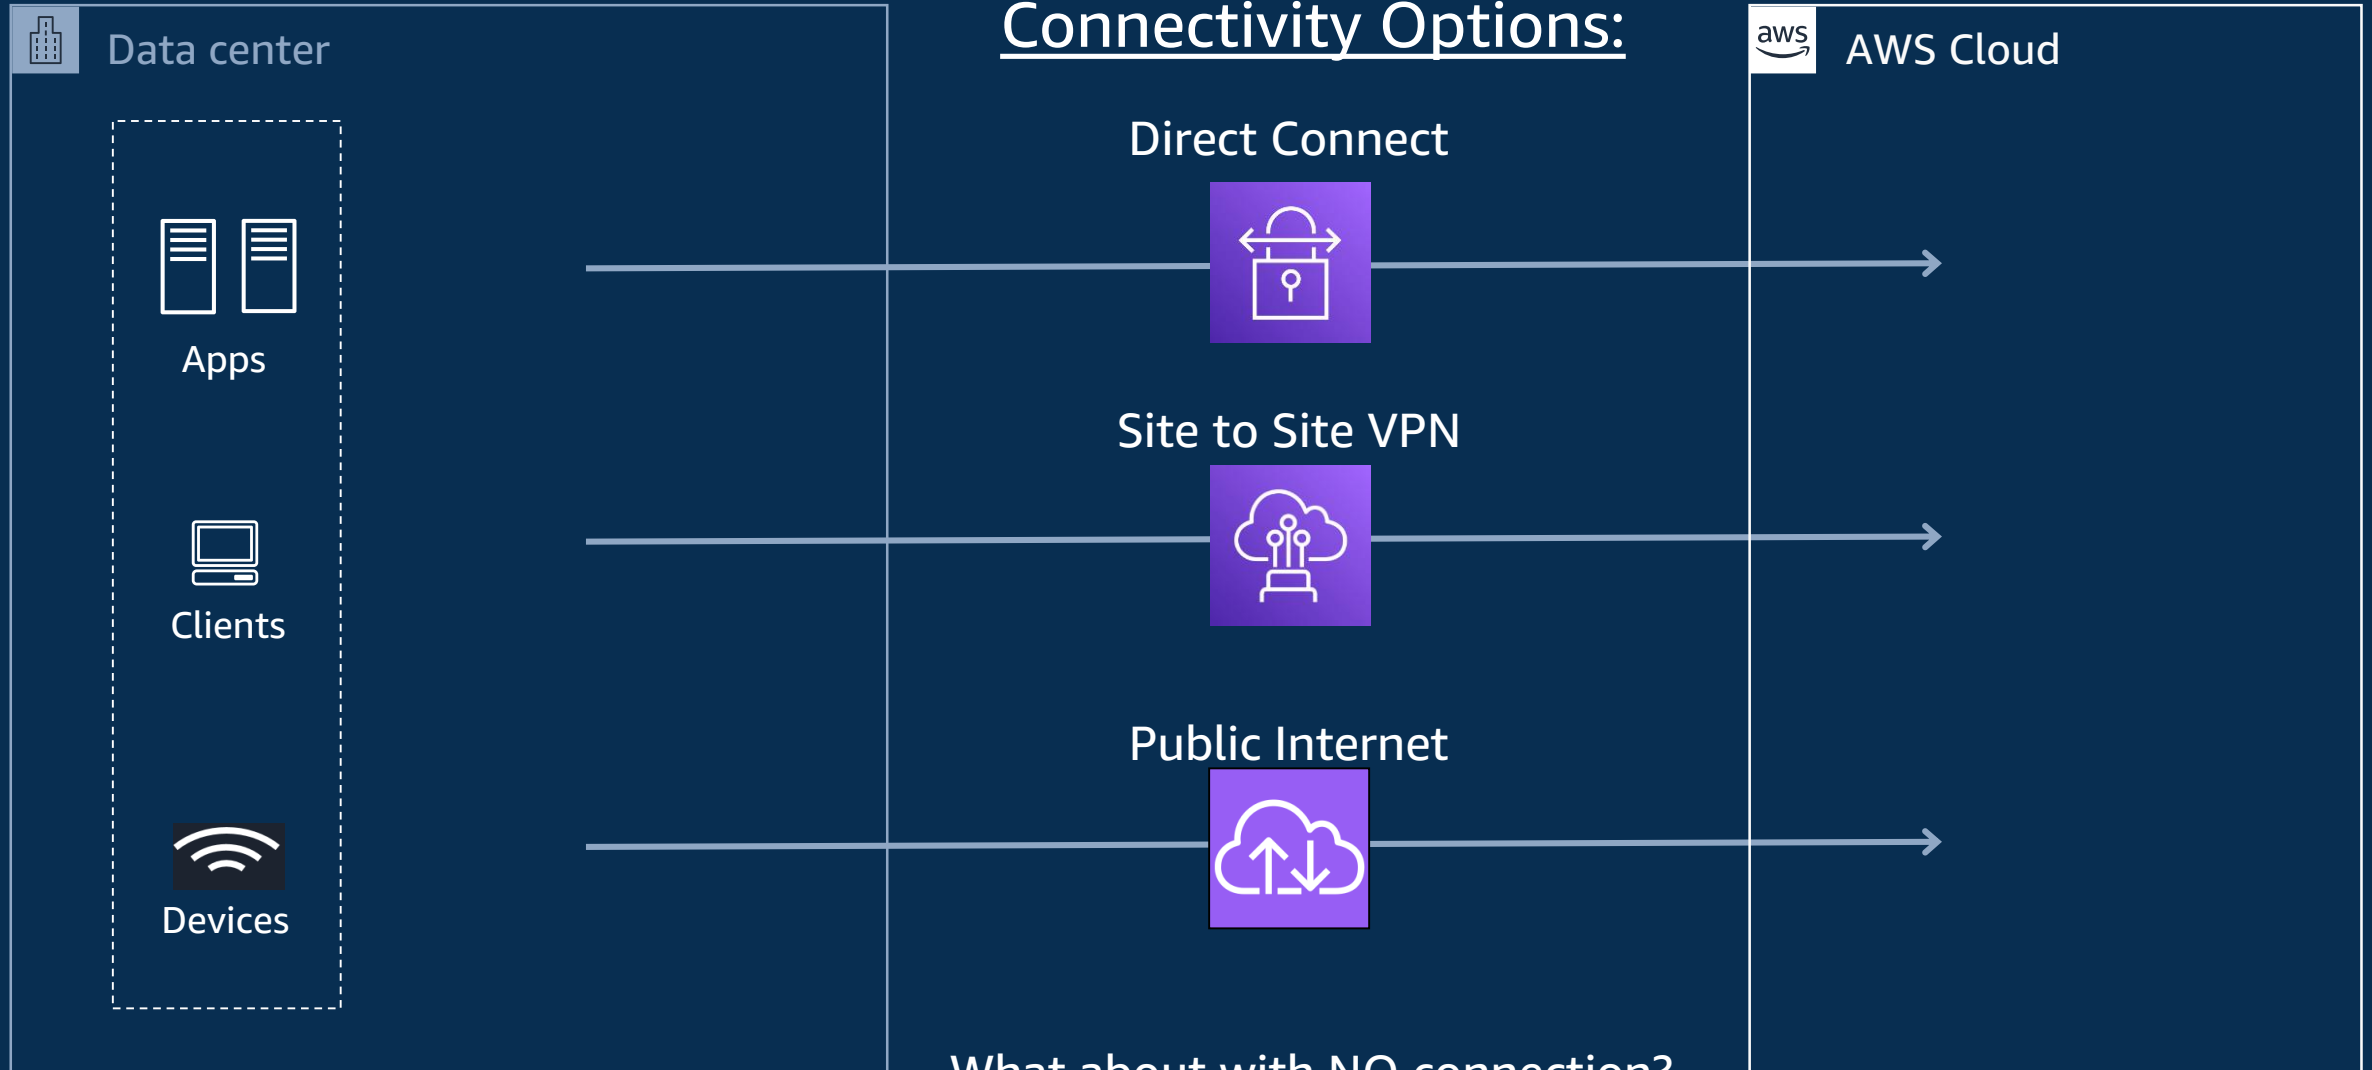
Connecting to AWS



Connecting to AWS



Connecting to AWS



What about with NO connection?

AWS storage portfolio

Data protection & resiliency



AWS storage portfolio



AWS Snow
Family



AWS Storage
Gateway




AWS DataSync

AWS DataSync



Data transfers are challenging



- Building and deploying scripts
- Encrypting and verifying data
- Recovering from errors
- Handling network availability
- Assuring performance

What is AWS DataSync?

Online data migration service
that simplifies, automates, and accelerates
copying file and object data to and from AWS storage



Fast data transfer

- Highly optimized, parallel network transfer (up to 100 TB/day)
- Transfers only incremental changes



Easy to use

- Schedule transfers
- Throttle bandwidth
- Filter by file name patterns



Secure and reliable

- End-to-end encryption
- End-to-end data verification
- VPC endpoints with AWS PrivateLink



Fully managed

- Integrates with AWS management and monitoring services
- Direct transfer into AWS storage services



Cost effective

- \$0.0125/GB transferred
- No minimums

AWS DataSync use cases



Migrate
application data
to AWS



Archive
on-premises
storage to free
up capacity



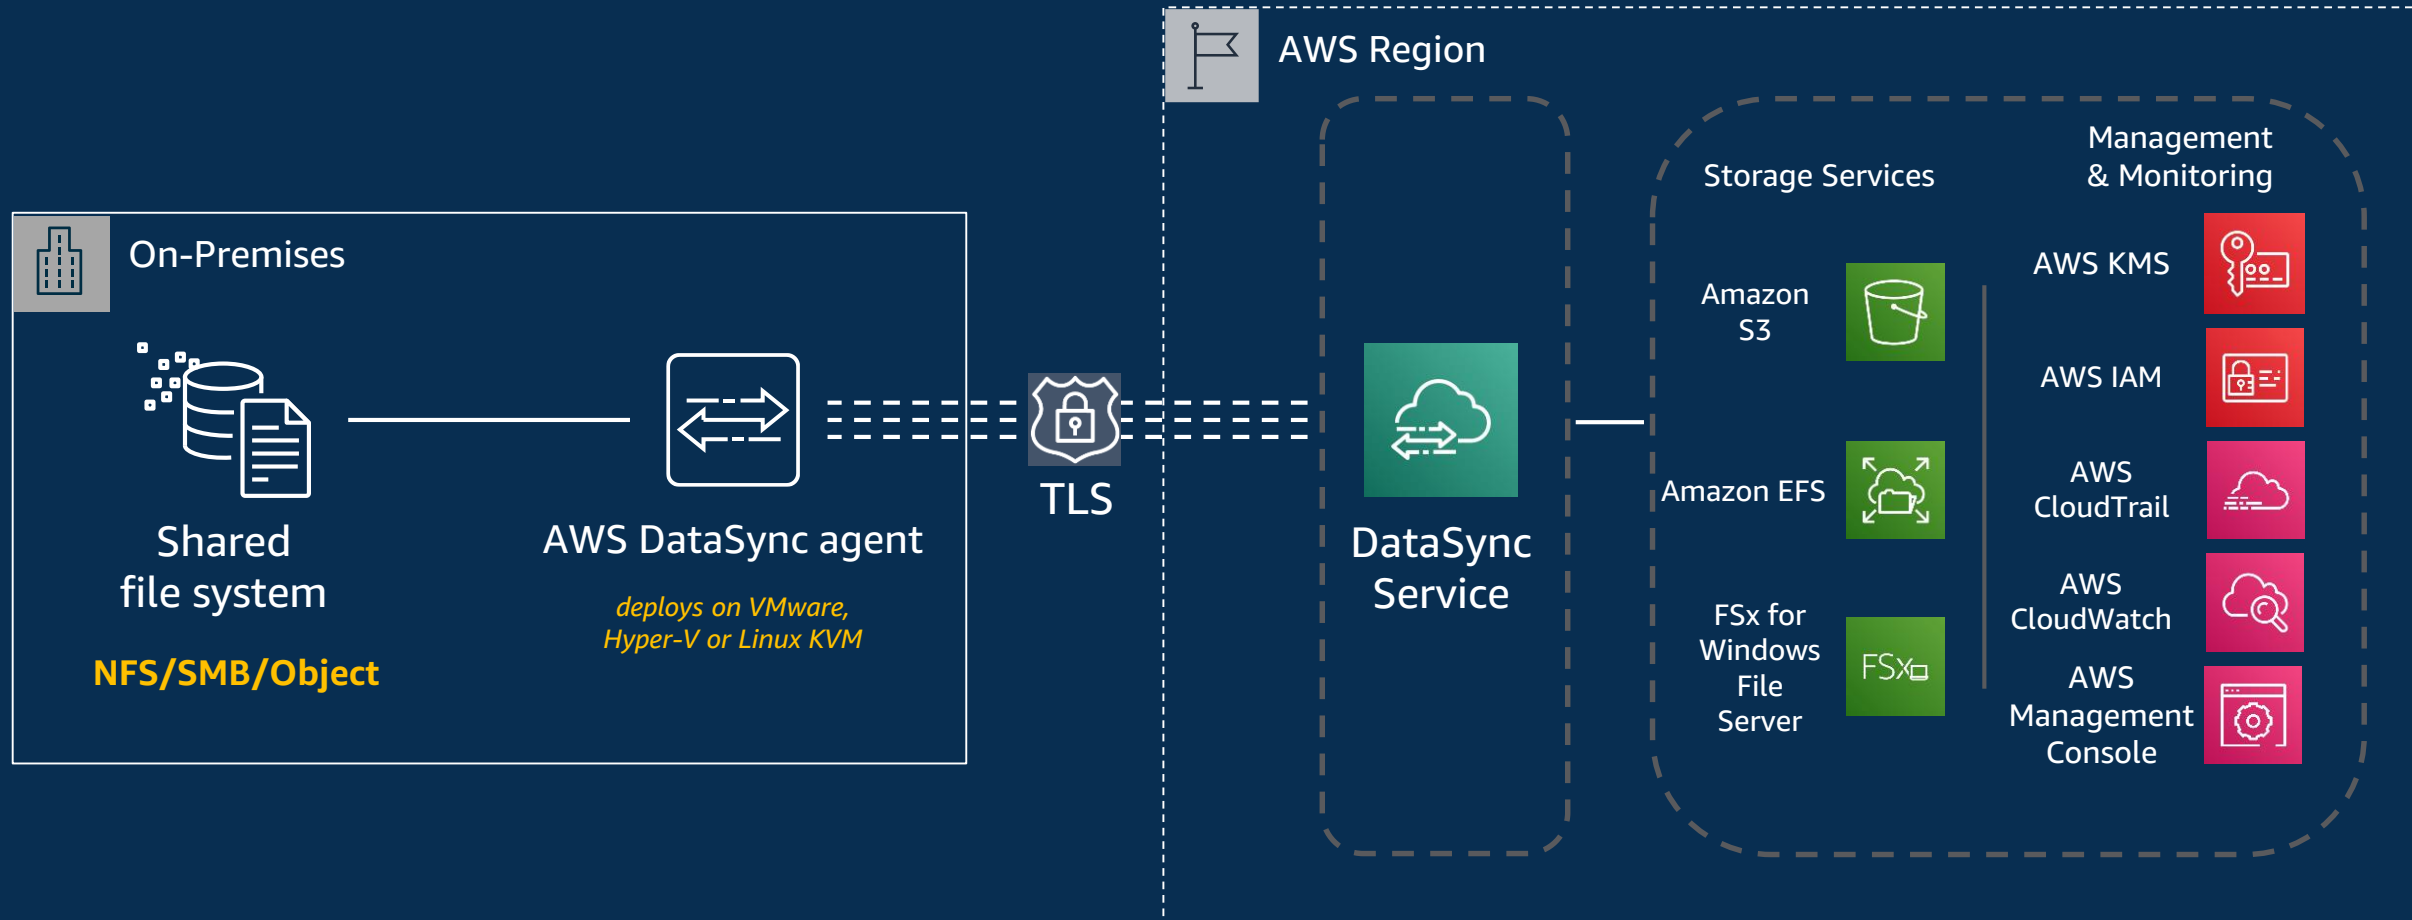
Replicate
for data protection
and recovery



Move
data for timely
in-cloud processing

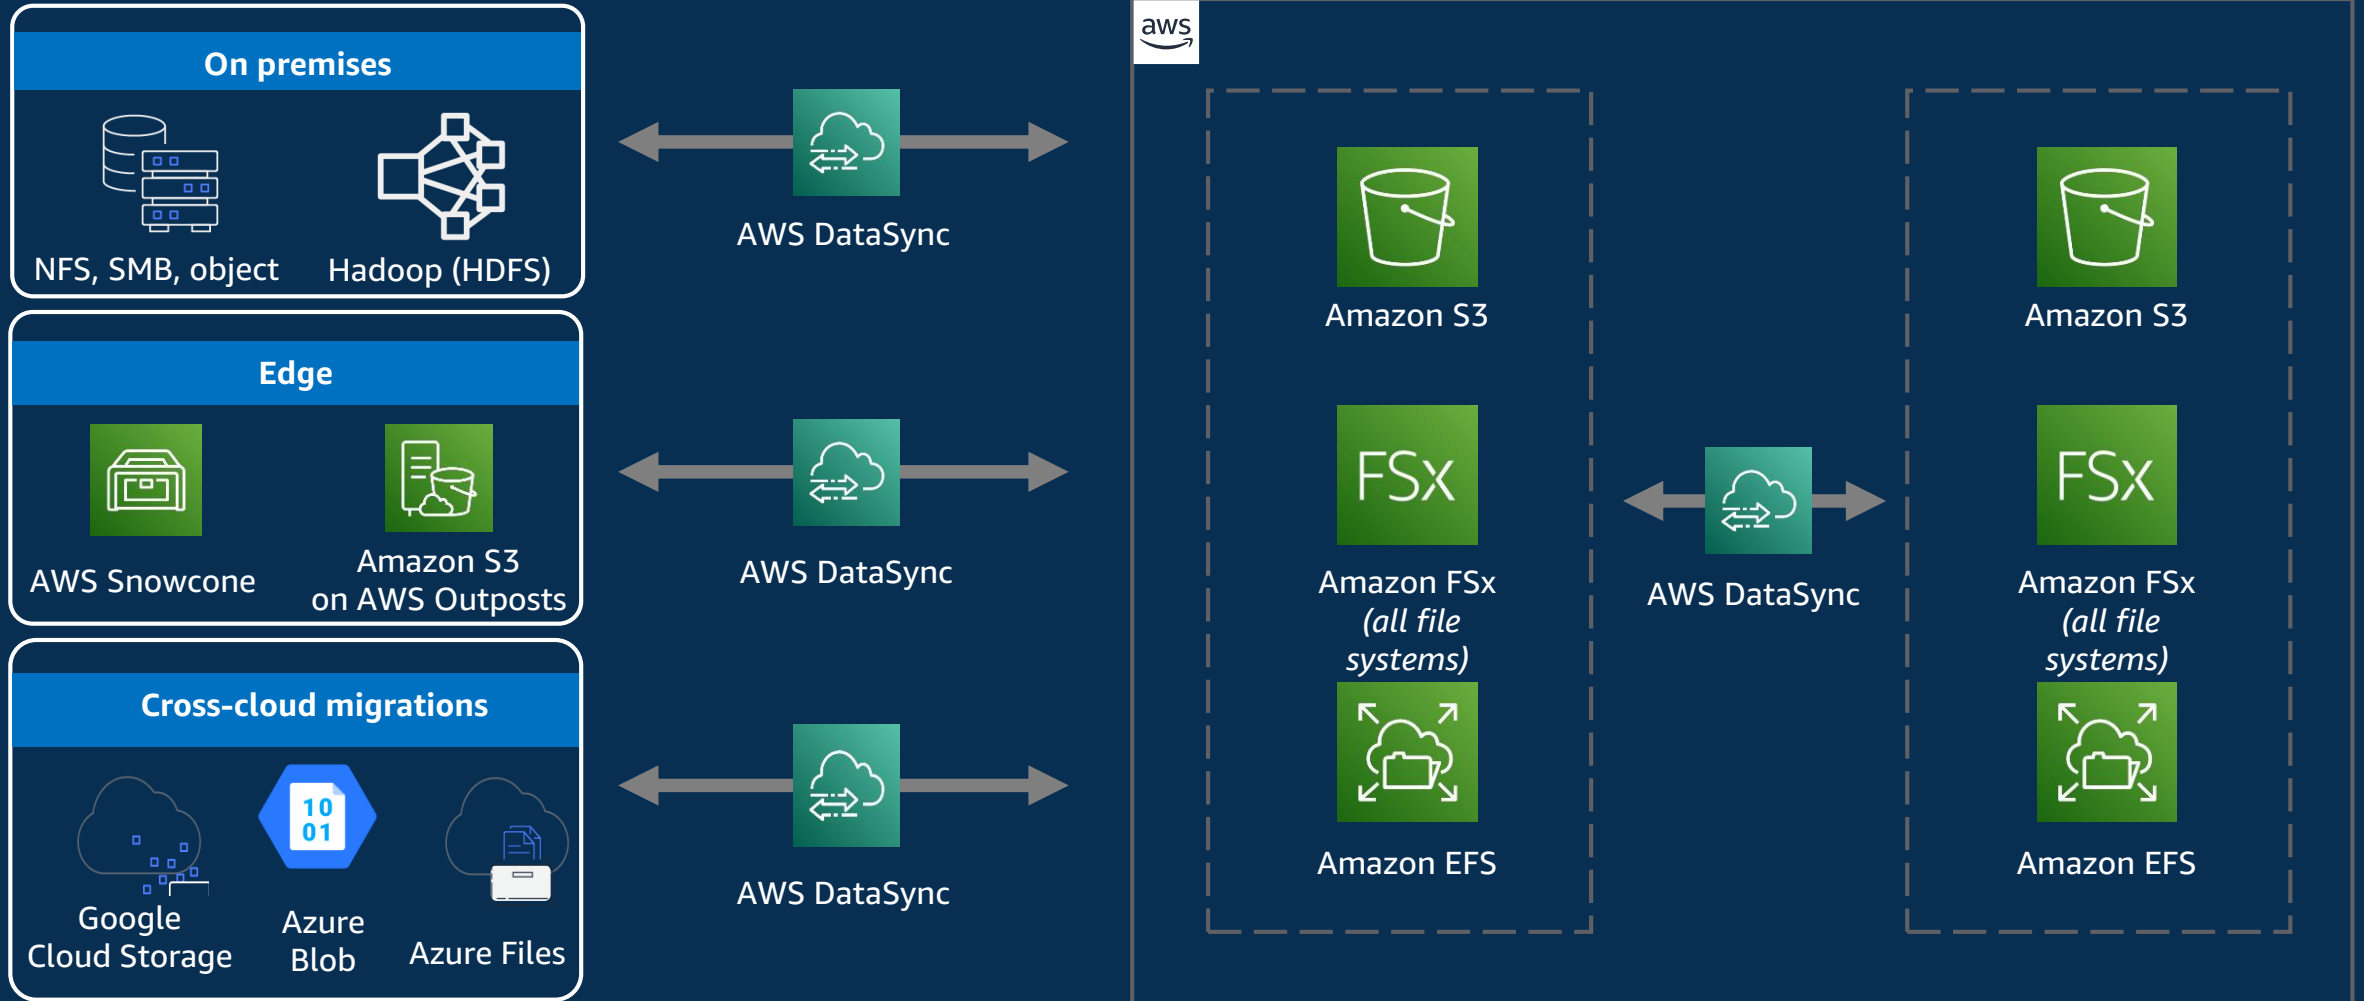
AWS DataSync architecture

SIMPLIFIES, AUTOMATES, AND ACCELERATES DATA TRANSFER TO OR FROM AWS

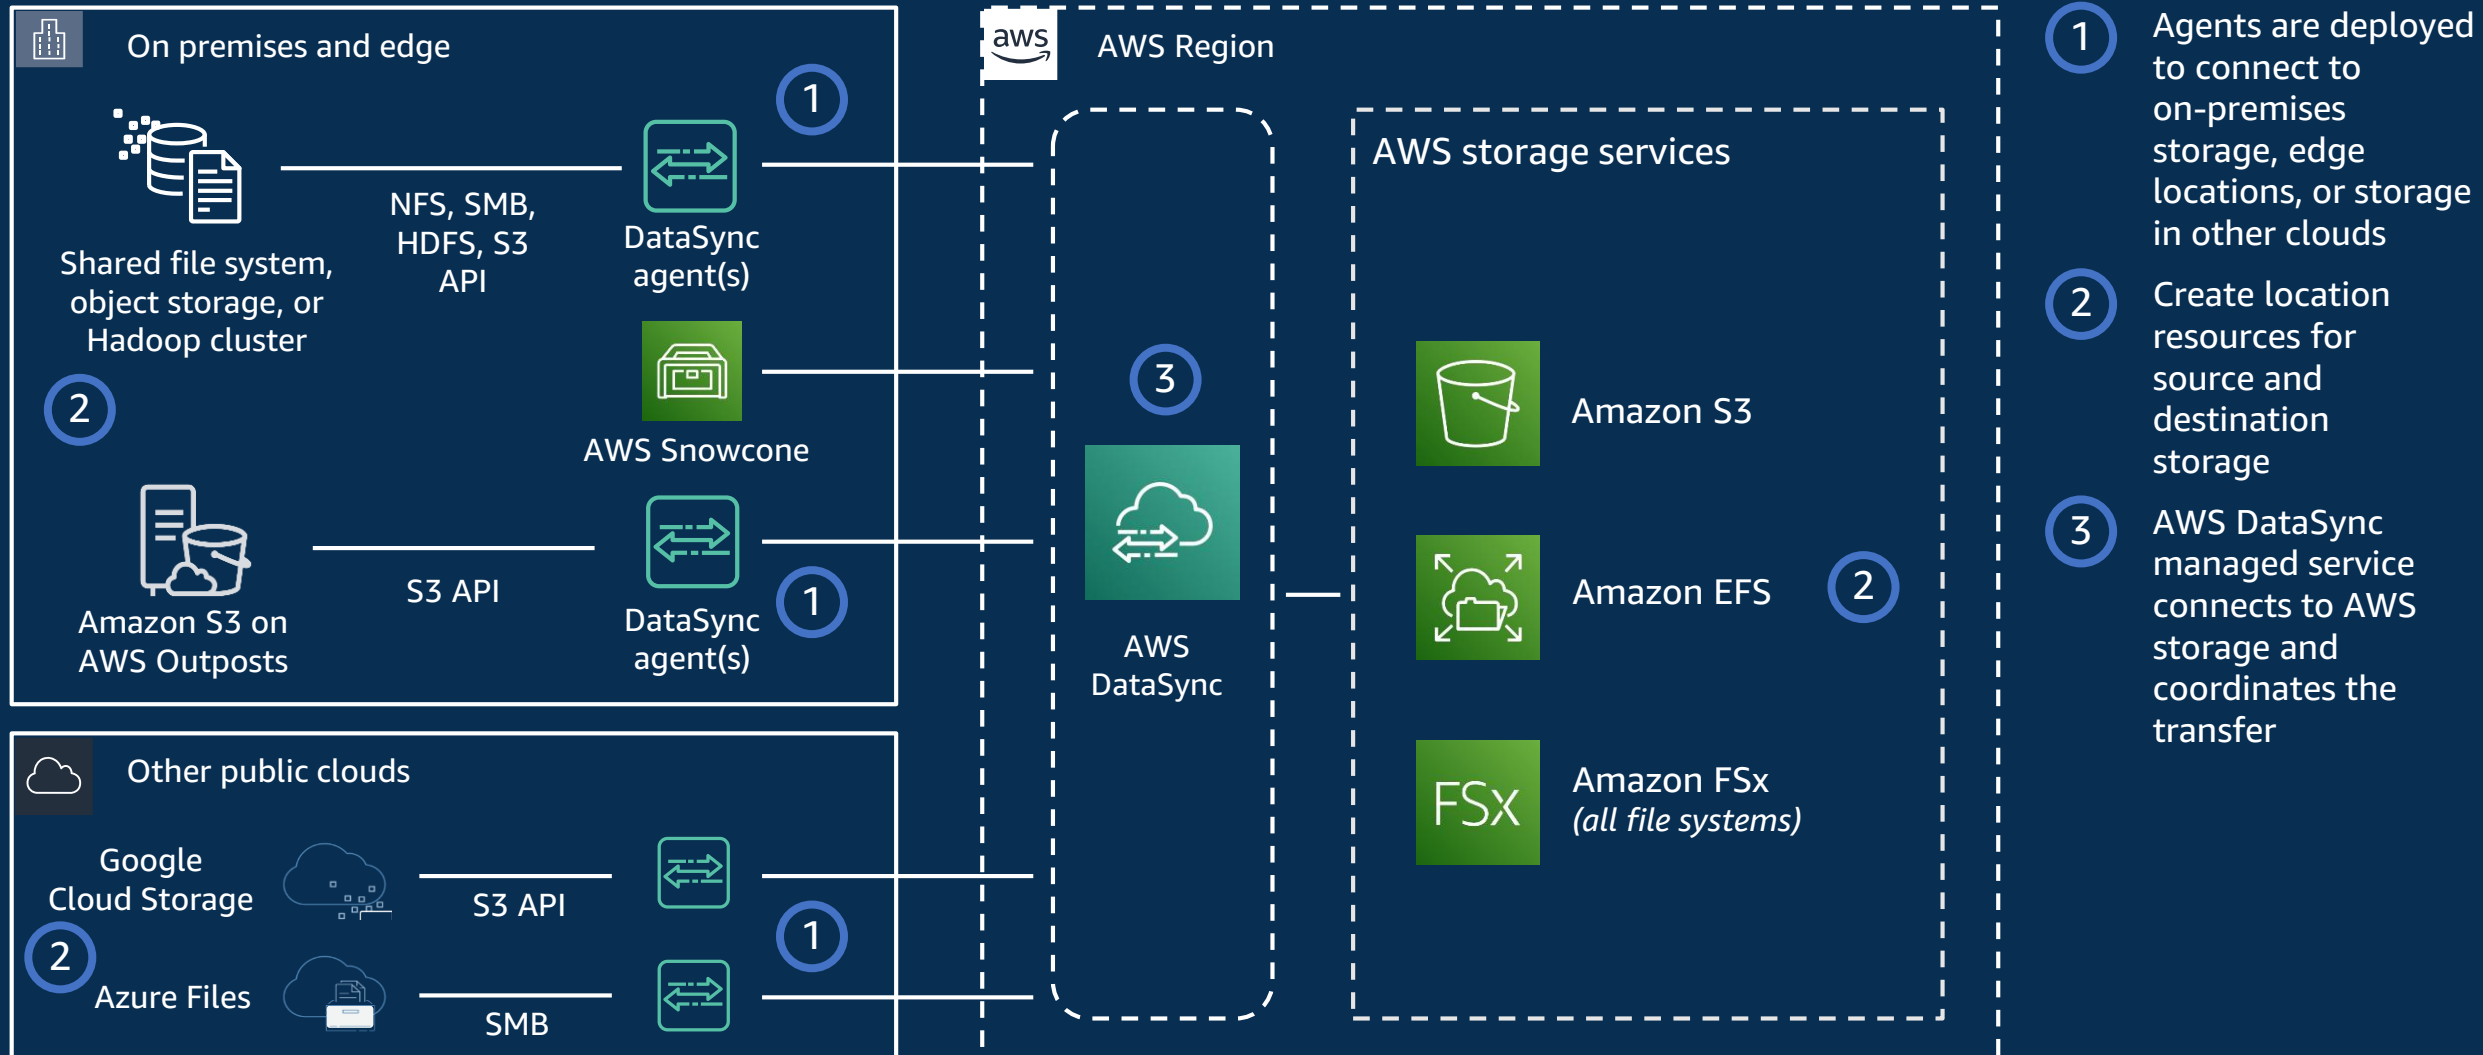


What can you do with AWS DataSync?

QUICKLY, SIMPLY, AND SECURELY MIGRATE YOUR DATA TO AND FROM AWS STORAGE SERVICES



AWS DataSync: How it works



AWS Storage Gateway



The hybrid cloud model

You have **on-premises** data and applications . . .



Hybrid cloud storage challenges



On premises

In the cloud



I am running out of on-premises storage capacity . . .

. . . and need to use scalable, durable, and secure cloud storage

I want to run my existing applications without change (files, backups, archive) . . .

. . . but leverage the infinite scale of the cloud

I need low-latency, on-premises access to data in the cloud . . .

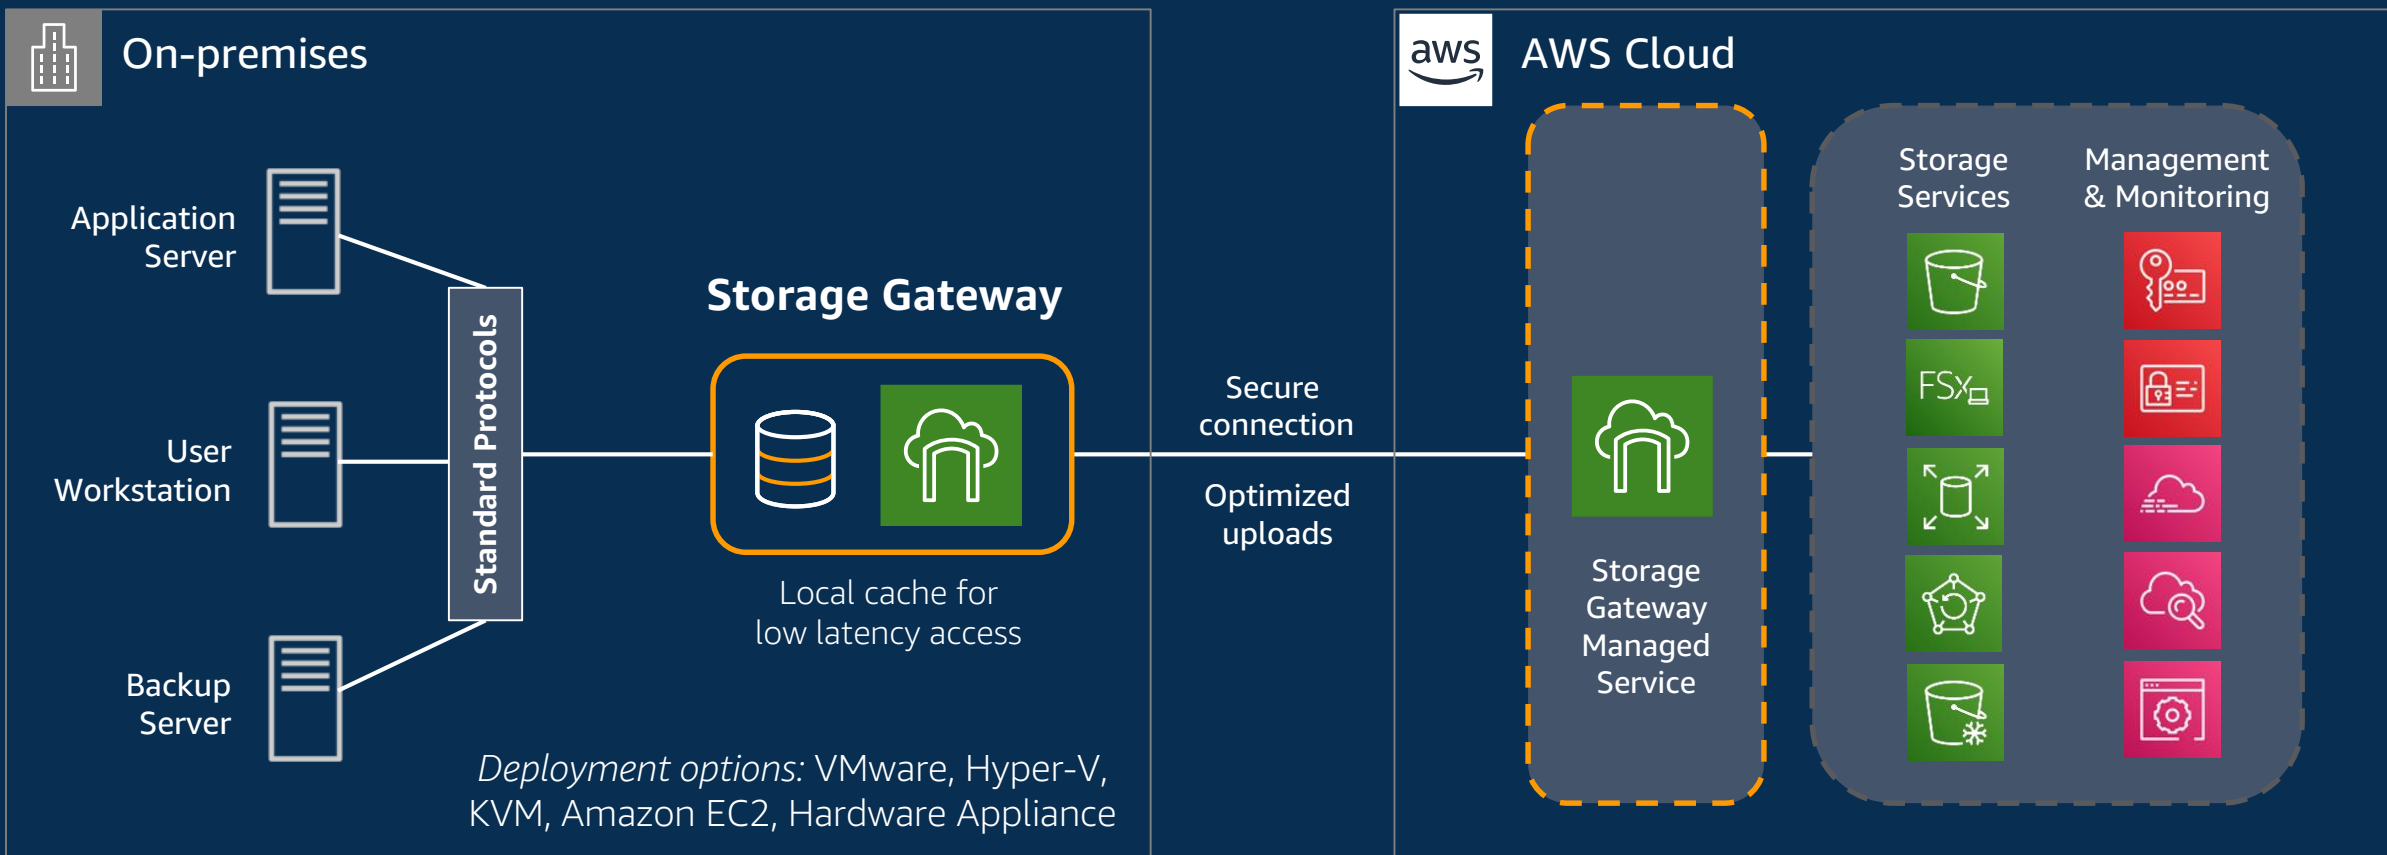
. . . while accessing the power of cloud compute and analytics

Compliance requires that I integrate with enterprise security and management tools . . .

. . . but I want to manage and monitor from a single pane of glass

Storage Gateway architecture

PROVIDES ON-PREMISES ACCESS TO VIRTUALLY UNLIMITED CLOUD STORAGE



AWS Storage Gateway

PROVIDING CLOUD-BACKED FILE SHARES, VIRTUAL TAPES, AND BLOCK VOLUMES



Amazon S3 File Gateway

Native file access, SMB or NFS, to Amazon S3 for **backups**, **archives**, and ingest for **data lakes**



Amazon FSx File Gateway

Replace on-premises Windows file shares with efficient, low-latency Amazon FSx for Windows File Server based **group file shares** and **home directories**



Tape Gateway

Replace **physical tape infrastructure**, leveraging Amazon S3 archive tiers for long-term retention




Volume Gateway

Block storage volumes with snapshots, AWS Backup integration, and cloud recovery

Storage Gateway benefits

PROVIDING ON-PREMISES APPLICATION ACCESS TO VIRTUALLY UNLIMITED CLOUD STORAGE



Standard protocols

No changes to existing apps



Local caching

Low-latency access to frequently used data



Optimized data transfers

Minimized network traffic



Provides security and supports compliance

FIPS, HIPAA, PCI, SOC*, ISO**, and encryption



Cost-effective

Pay-as-you-go pricing




AWS integrated

Management, monitoring, and in-cloud workloads

*SOC (1, 2, 3) **ISO (9001, 27001, 27017, 27018)

Demo

Amazon S3 File Gateway

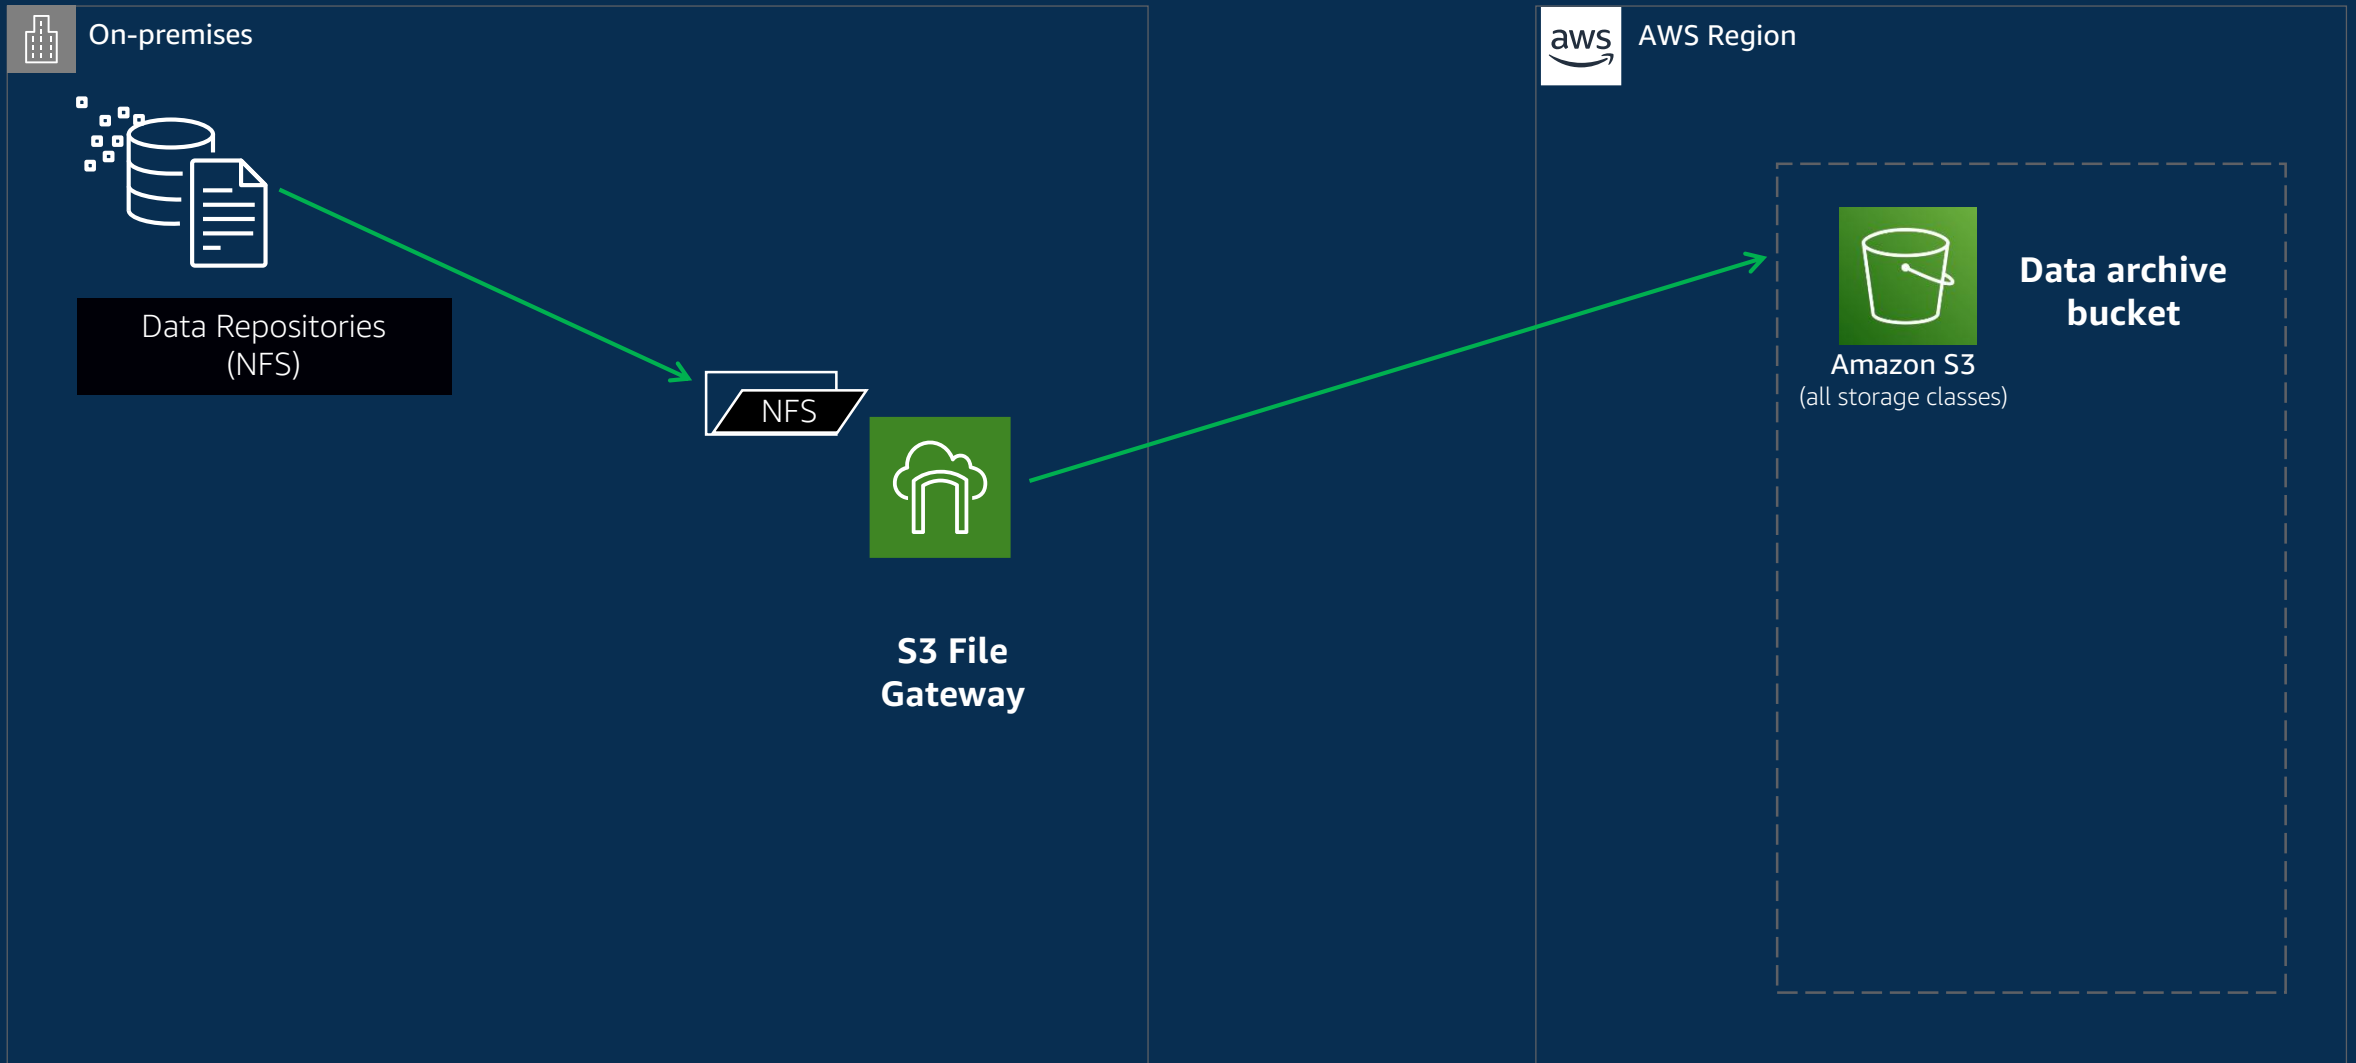


Demo Scenarios – Sophia and Joe

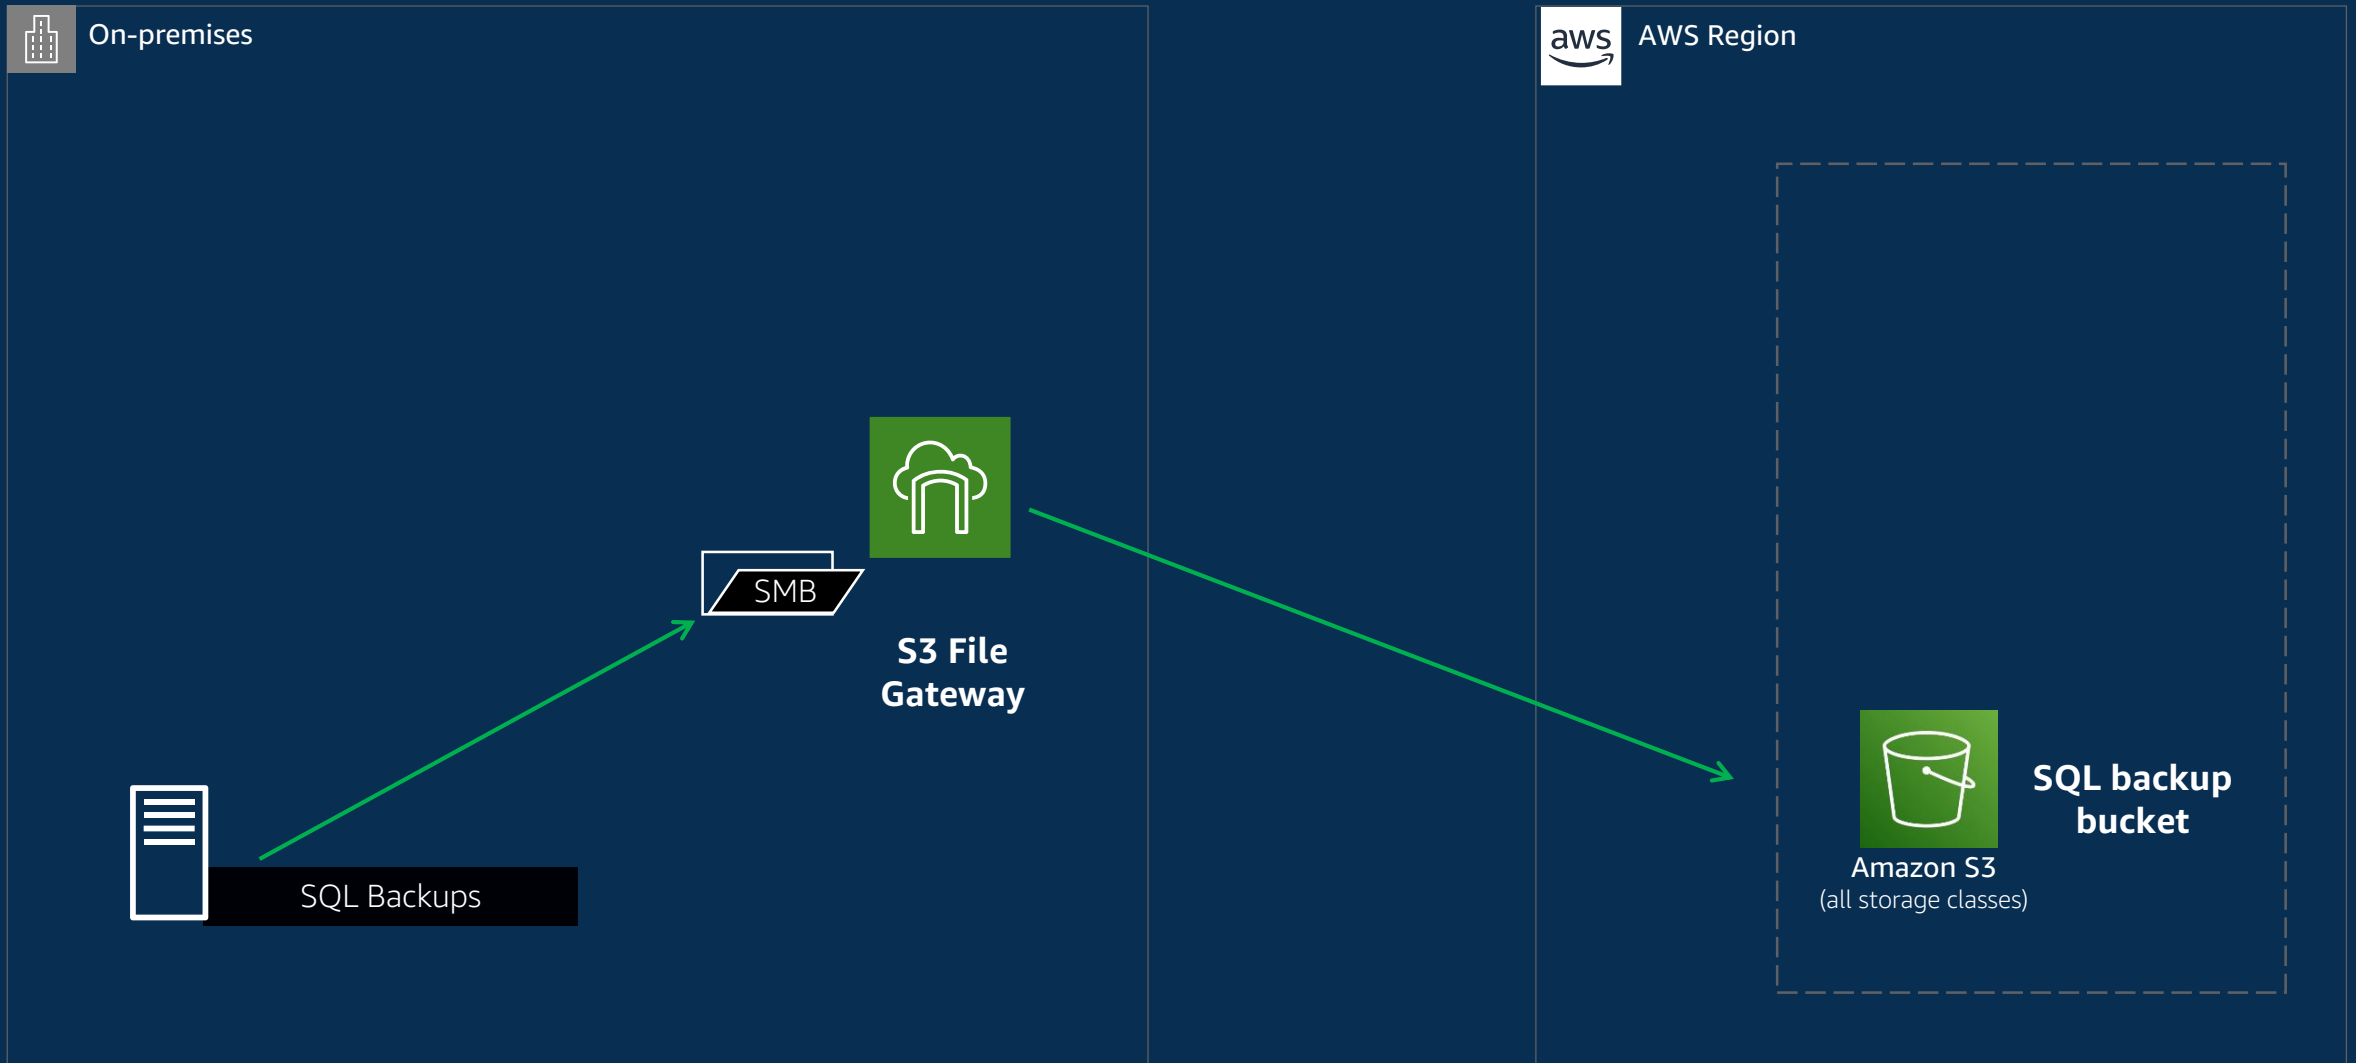
Sophia has large data repositories stored on NAS Storage presented via **NFS** which need to be **archived and migrated** into AWS, and **still available via NFS** on-premises.

Joe manages Backups and has run out of on-premises backup storage – and has realized the **SQL backups** are taking up a large portion of their expensive primary storage.

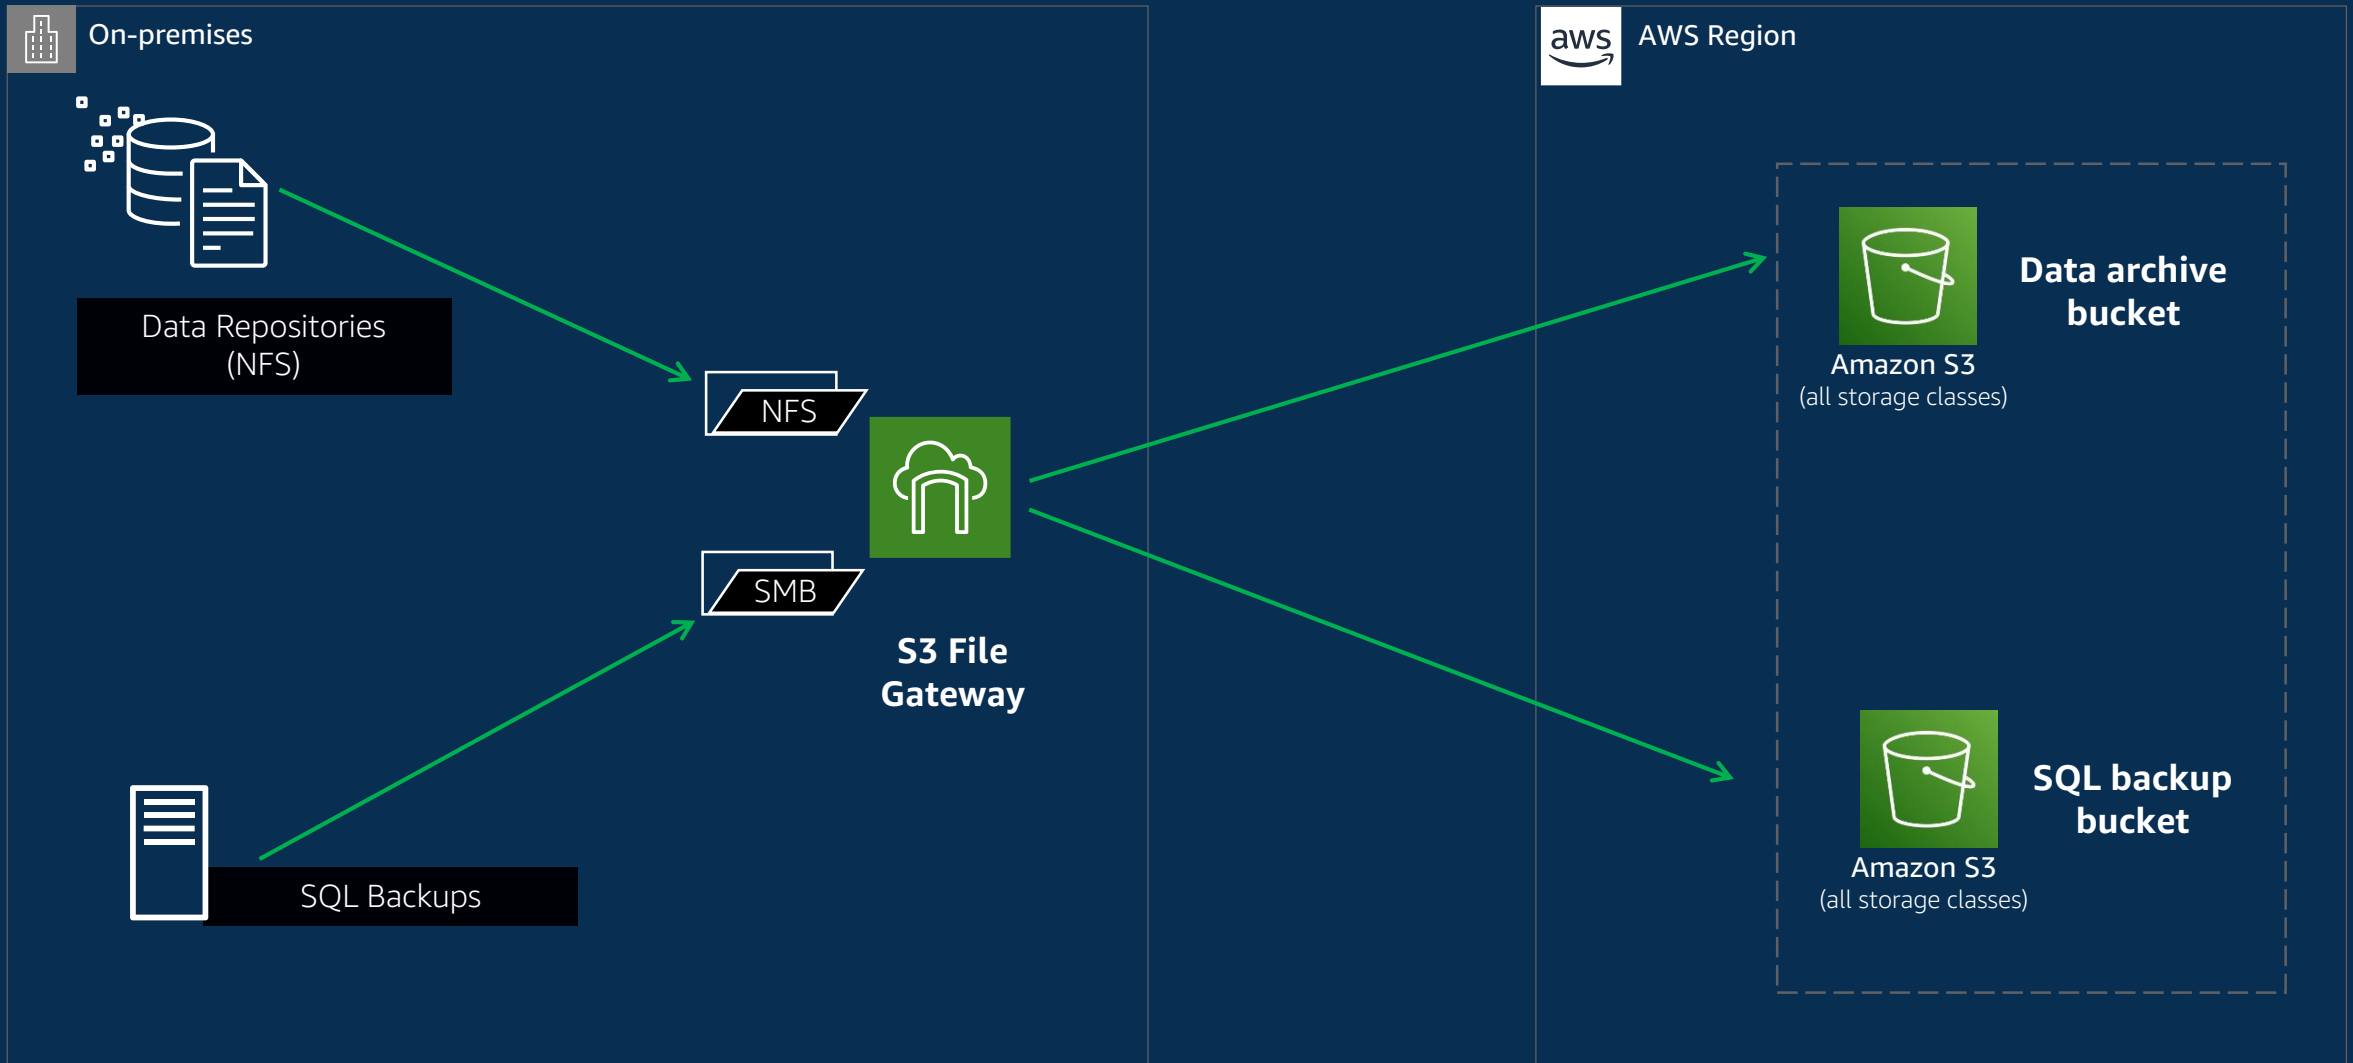
Demo set up 1 – S3 File Gateway (NFS)



Demo set up 2 – S3 File Gateway (SMB)



Demo set up – S3 File Gateway (NFS and SMB)



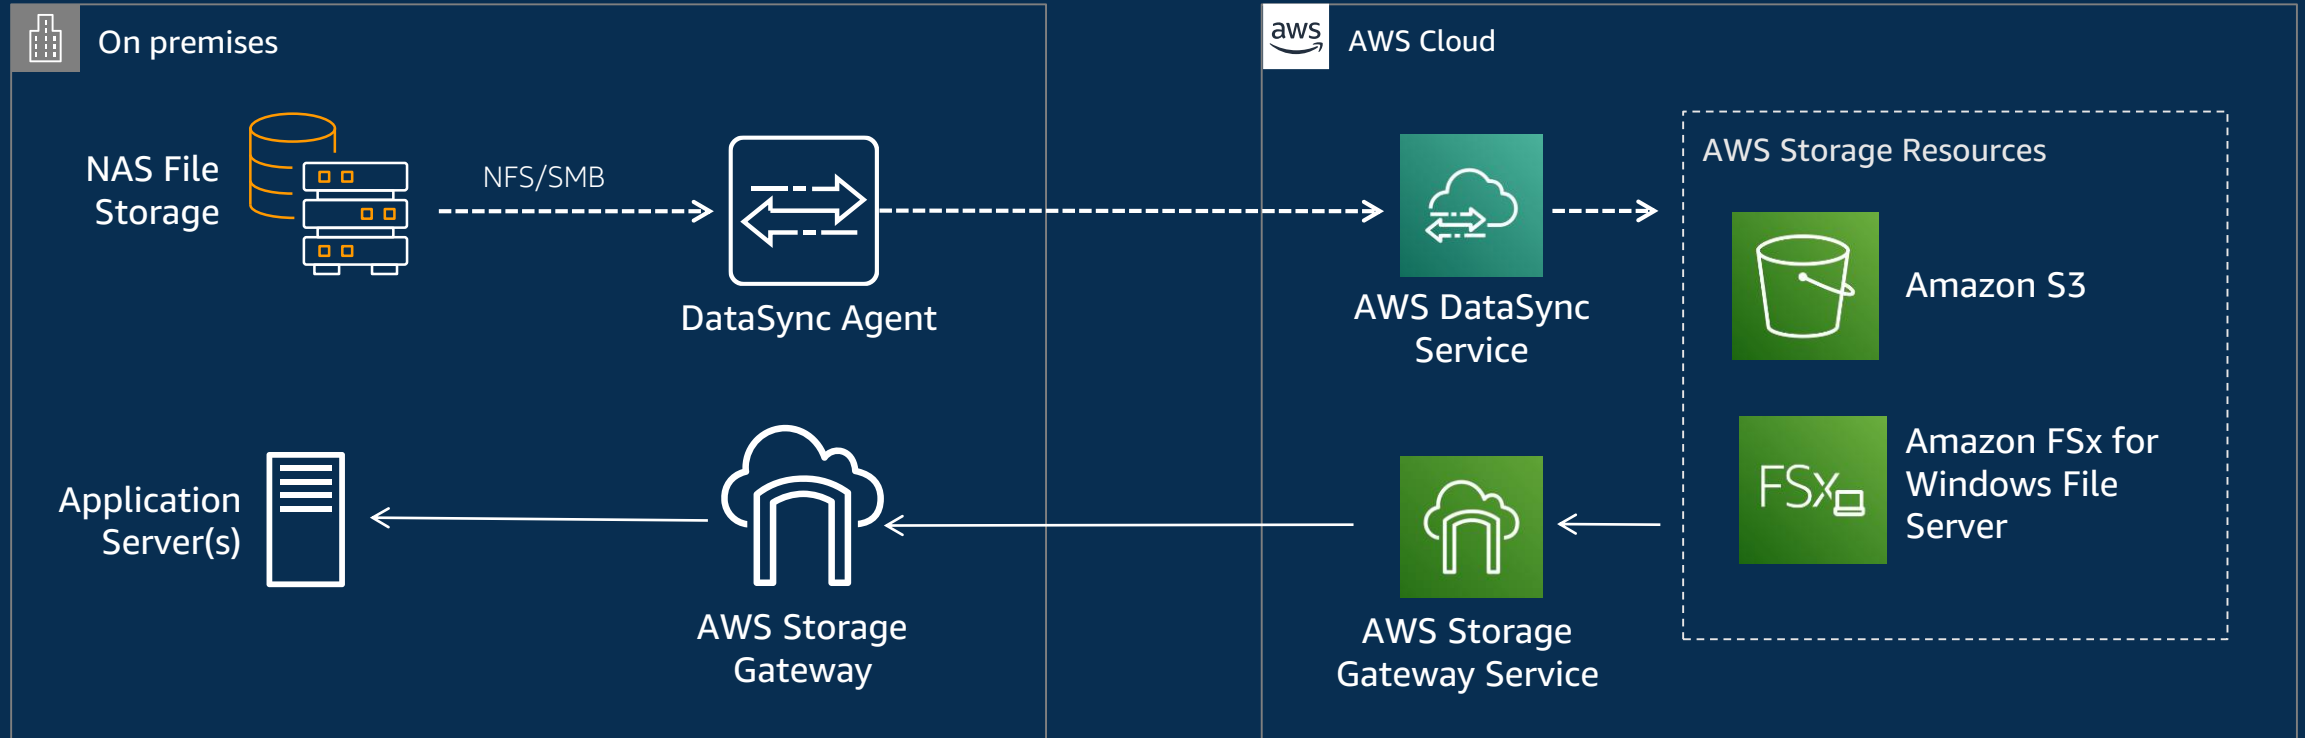
Demo recap

- Amazon S3 File Gateway was used to provide **low latency access** to an **applications** for on-premises access.
- S3 File Gateway was used to **transfer / migrate** data from **different sources** on premises (NFS and SMB) to Amazon S3
- A **seamless way to connect to the cloud** (via SMB / NFS) in order to store **backup images** as durable objects in Amazon S3 cloud storage.
- Objects written **through the S3 File Gateway** can be **accessed directly in S3**.
- A single gateway can connect to **multiple buckets**.

Common hybrid pattern

Archive data to the cloud

MOVE CLOUD DATA FROM EXPENSIVE ON-PREMISES STORAGE SYSTEMS TO DURABLE, SECURE LONG-TERM STORAGE



FORMULA 1 uses AWS DataSync to replicate video archive to Amazon S3



Challenge

- 400 TB video archive that needs to be backed up regularly while maintaining high accessibility

Solution

- Using AWS DataSync, FORMULA 1 copies up to 4 TB per day to Amazon S3.
- FORMULA 1 uses AWS Storage Gateway to provide standard file access to data in Amazon S3 with a local cache for low latency access.

Results

- AWS DataSync transfers data quickly and securely to Amazon S3.
- Data is stored durably and reliably in Amazon S3, and all data is encrypted at rest.
- AWS Storage Gateway gives users on-premises access to the video archive in Amazon S3, providing fast recovery from failure scenarios.

<https://aws.amazon.com/blogs/storage/formula-1-transfers-race-car-data-into-the-aws-cloud-using-aws-datasync/>



“AWS DataSync synced data from our on-premises file server to Amazon S3 at an average of 4 TB per day, and did so with encryption during transfer and at rest.”

Martynas Juras

Cloud Engineer, FORMULA 1

AWS Snow Family



What is AWS Snow Family?

Move petabytes of data to and from AWS,
or process data at the edge.



AWS
Snowcone



AWS
Snowcone SSD



AWS
Snowball

<1TB



10s PB


AWS Snow Family data migration workflow



AWS Snow Family management console
Qualify SBE V3S then
Create SBE V3S job



Use AWS OpsHub or
AWS CLI to manage
AWS Snow Family
devices



Courier transport
service



Secure erasure
using NIST 800-88 standards

Device
shipped

Data copied
1-N* days

Device arrives
at AWS

1-3 days

1-2+ days

1-2+ days

1-N days

Job
created &
provisioned

Device
arrives

Device sent
to AWS

Job
Complete



Courier transport
service



Copy data via >1GB/s
local network



Your data begins ingesting
into your Amazon S3
bucket



Recap

- Hybrid cloud storage services
- AWS DataSync
- AWS Storage Gateway
- AWS Snow Family devices
- Demo of using AWS Storage Gateway for NFS and SMB on-premises

Additional Resources

- AWS DataSync : <https://aws.amazon.com/datasync>
- AWS Storage Gateway : <https://aws.amazon.com/storagegateway>
- AWS Snow Family : <https://aws.amazon.com/snow>
- AWS Storage : <https://aws.amazon.com/products/storage/>

Get hands-on experience with the AWS online data migration workshop

<https://github.com/aws-samples/aws-online-data-migration-workshop>

AWS TRAINING & CERTIFICATION

Access 600+ free digital courses with AWS Skill Builder

Focus on the cloud skills and services that are most relevant to you across 30+ AWS solutions, including digital self-paced learning plans and ramp-up guides.

LEARN YOUR WAY [EXPLORE.SKILLBUILDER.AWS](https://explore.skillbuilder.aws) »



Validate your cloud expertise with an AWS Certification

Take the step towards earning an industry-recognised credential. Learn more about how to become an AWS Certified Cloud Practitioner, and AWS resources that can help you prepare.

ACCESS RESOURCES TO [PREPARE FOR YOUR EXAM](#) »



Thank you for attending AWS Builders Online Series

We hope you found it interesting! A kind reminder to **complete the survey**.
Let us know what you thought of today's event and how we can improve the event experience for you in the future.



aws-apj-marketing@amazon.com



twitter.com/AWSCloud



facebook.com/AmazonWebServices



youtube.com/user/AmazonWebServices



linkedin.com/company/amazon-web-services



twitch.tv/aws



Thank you!

Naim Mucaj

Senior Solutions Architect
Amazon Web Services



Key takeaways

AWS DataSync moves and copies data from on premises to AWS and between AWS storage services

AWS DataSync moves data online using your network bandwidth or a VPC for added security

AWS DataSync works well with the AWS Snow Family of devices; DataSync is built into AWS Snowcone for edge use cases

Get started with AWS DataSync by visiting aws.amazon.com/datasync